



UNIVERSITY OF PISA

Master in Computer Engineering
Project for Electronic Systems

Perceptron

Author:

Jacopo Carlon

587412

Academic year

2022/2023

Contents

1	Introduction	2
1.1	Specification	2
1.2	Circuit Employments	4
2	Architecture	5
2.1	Sum-Product	5
2.2	Activation Function	5
2.3	Other optimizations	6
2.4	Timing and inputs	7
3	VHDL code	8
3.1	Entity List	8
3.2	Perceptron : code	8
3.3	Test Bench : code	14
4	Verification and Testing	23
4.1	Test-Bench output example	23
4.2	Expanding the Test-Bench	24
5	Synthesis and Implementation	25
5.1	Vivado design flow	25
5.2	RTL	25
5.3	RTL Elaboration	25
5.4	Synthesis and Implementation	26
5.5	Troubleshooting and warnings	27
6	Vivado Results	28
6.1	Critical Path	28
6.2	Timing Report	29
6.3	Utilization Report	30
6.4	Power Report	30
7	Final Considerations	31
7.1	VIVADO results analysis	31

1 Introduction

1.1 Specification

It is requested to design a perceptron which takes $N_{IN} = 10$ inputs x_n represented with $b_x = 8$ bits. The network will carry out products between x_n and generic coefficients w_n , adding to the result a bias b . w_n and b are represented with $b_w = 9$ bits.

x_n , w_n and b shall be considered in range $[-1, 1]$ using fixed point arithmetic.

This means that the inputs will have 2 bits for the integer part and 6 (for x_n) or 7 (for w_n and b) bits for the fractionary part.

The output y of the system shall be represented with $b_y = 16$ bits, truncating the least significant bits.




Figure 1: Perceptron

The activation function of the perceptron shall be the one represented in the next figure:




Figure 2: Activation Function

In the design used for Vivado implementation there will be registers catching all the inputs at *rising_edge(clk)*, the operations will be evaluated in a combinatory manner, and the output will be assigned to the *output – register* at the following *rising_edge(clk)*.




Figure 3: Perceptron Interface

1.2 Circuit Employments

A Perceptron, in Deep Learning, is a neural network link that contains computations to track features and use Artificial Intelligence in the input data. These artificial neurons are obtained using simple logic gates with binary outputs.

The perceptron is equivalent to an artificial neuron, which invokes the mathematical function and has node, input, weights, and output equivalent to the cell nucleus, dendrites, synapse, and axon, respectively, compared to a biological neuron.

In 1957 Frank Rosenblatt proposed a Perceptron learning rule based on the original MCP neuron.

A Perceptron is an algorithm for supervised learning of binary classifiers.

This algorithm enables neurons to learn and process elements in the training set one at a time.

This perceptron (and a combination of) can be used as an accelerator to improve neural network performances (e.g. during training phase).




Figure 4: Perceptron Idea

2 Architecture

This implementation of the perceptron is obtained by having one main block process, used to assign default values to all the registers at the *asynchronous reset*, and to assign input values to the input-registers at *rising_edge(clk)*. This main block also receives from the combinatory logic the elaborated y value in 16 bits, and assigns it to the output register at *rising_edge(clk)*.

Other than the main logic, there is a big combinatory part used to implement the product of x_j , w_j for the j inputs, adding them together and adding the *bias* b . From this prod-sum, the final output is generated by further combinatory logic, which implements the activation-function.

2.1 Sum-Product

Since we have inputs x_i in 8 bits (2 int & 6 fract) and weights w_i (2 int & 7 fract), and since they are to be considered *signed* (they range in $[-1, 1]$), each product result should be *signed* in 17 bits (4 int & 13 fract).

We also need to add -once- the bias (2 int & 7 fract), which would need to be expanded -with sign- to 18 bits at the least. The overall result "*sp_res*" of these 11 sums would be in 4 more bits to a total of 21 bits (8 int & 13 fract).

However, this size is not necessary! We know that all inputs are in $[-1,1]$, therefore all products will be in $[-1,1]$ and the sum will be in $[-11, +11]$. Since we do not want to lose accuracy, and since we need to keep in mind that these are *signed* values, we can represent *sp_res* in just 18 bits, 5 for the integer part and 13 for the fractionary part.

This is because we know that the $[-11, +11]$ range in binary will be like [10101 , 01011]

In order to implement the operations so far described, we expand-with-sign all products to 18 bits (5 int & 13 fract) and sum them without any fear of overflows.

2.2 Activation Function

The function is to be read as follows : if *sp_res* is less than -2, return 0;

- if *sp_res* is less than -2, return 0;
- if *sp_res* is greater than 2, return 1;
- else : map the values from $[-2,2]$ to y $[0,1]$ on a straight line.

It must be noted that, since the output will be in 16 bits, it will have 3 of integer part and 13 bits of fractionary part (we are required to truncate -during the division-). The output will therefore be in one of three forms : 00-0-.00...00, 00-1-.00...00 or 00-0-.XX...XX, which means the output will always be positive !

To implement the comparison we simply do a comparison with sign.

To implement the tilted-line part, however, we need some more considerations.

We could implement the function $y = x/2 + 1/2$, or the function $y = (x+2)/4$. Considering that we enter the -else- block only if *sp_res* is in [-2,2], in order to get rid of the heavy *signed* property, I will firstly add 2, (represented in 5 int & 13 fract). By doing so I'll obtain positive number "*sp_plus_two*" (5 int & 13 fract) which will be in range [0, 4] and therefore will be in the form **00XXX.fract_13**.

The following operation of division by 4 could easily be implemented by a double shift to the right (padding the left with 0 since we have positive numbers is not a problem).

This would give us a value in the form *SS00X.XXfract_11* (5 int & 13 fract) with *S*=0, and we would need to later reduce it to **00X.XXfract_11** (3 int & 13 fract) for the output in 16 bits.

However, we can optimize this even more!

The final result can easily be achieved by simply extracting from *sp_plus_two* only 16 bits starting from the most significant bit (the leftmost one).

[Note that this elegant solution is so easily implementable because we know that the value is positive after adding 2!

This is the true reason for the above choice of function to implement.]

2.3 Other optimizations

After the first synthesis I noted that Vivado was implementing the sum of partial-products as a chain sum (cascading sum) which inevitably was causing delays in the generation of the final value.

In order to avoid such problems, I have manually rewritten the sum operations in the **structure of a balanced binary tree**.

This approach needed more signals (for the partial sum-prods), but reduced the sum portion of the combinatory part from 10 to **only 4 logic levels**, therefore allowing to decrease the minimum clock period (==increase the maximum clock frequency).

This also uses exactly the same number of adders, so no costs were added to this re-organization!

2.4 Timing and inputs

Since the time theoretically required to generate the output from given inputs is only due to the combinatory logic part, the timing is formalized as follows: - @*rising_edge*(*clk*)[*t*] inputs[*j*] are saved, current output is stabilized and kept - output[*j*] is elaborated - @*rising_edge*(*clk*)[*t* + 1] inputs[*j* + 1] are saved, output[*j*] is stabilized and kept

This means that, as long as all inputs are "as expected" (in range [-1,1] in 2-bits-int & 6|7-bits-fract), the Perceptron will be **sensible to inputs only at rising_edge(clk)[t]** , and the output of these inputs is **assured to be correct (exactly and only) at the following rising_edge(clk)[t+1]** .

Should inputs not be in the correct range, then the behaviour is N.D.:

Depending on the value (and sign!!!) of *sp_res*, the output could exchange 0 for 1, or mistakenly fall in the "else".

3 VHDL code

In this section it is shown and commented the VHDL code used to design the Perceptron. The test bench (TB) source is reported here with a brief description, however its logic is analysed in the section dedicated to the verification approach.

3.1 Entity List

Due to the extreme simple nature of the component, only a single VHDL entity was implemented:

- Perceptron

3.2 Perceptron : code

The following code shows a main process sensible to (asynchronous) reset and clock. The main process has a number of registers corresponding to the 21 inputs , and a register for the output y. The architecture also has a considerable combinational logic part, which works as described above, in order to calculate all the products and sums, and in order to implement the activation function of the Perceptron.

```
1
2 library IEEE;
3 use IEEE.std_logic_1164.all;
4 use IEEE.numeric_std.all;
5 use IEEE.std_logic_unsigned.all;
6
7 entity Perceptron is
8     port(
9         clk          : in std_logic;
10        resetn       : in std_logic; -- active low
11
12        -- input values xi 8 bit, fpv, in [-1, 1]
13        x0:          in std_logic_vector(8 - 1 downto 0);
14        x1:          in std_logic_vector(8 - 1 downto 0);
15        x2:          in std_logic_vector(8 - 1 downto 0);
16        x3:          in std_logic_vector(8 - 1 downto 0);
17        x4:          in std_logic_vector(8 - 1 downto 0);
18        x5:          in std_logic_vector(8 - 1 downto 0);
19        x6:          in std_logic_vector(8 - 1 downto 0);
20        x7:          in std_logic_vector(8 - 1 downto 0);
21        x8:          in std_logic_vector(8 - 1 downto 0);
22        x9:          in std_logic_vector(8 - 1 downto 0);
```

```

23     -- input weights wi 9 bit, fpv, [-1, 1]
24     w0:      in   std_logic_vector(9 - 1 downto 0);
25     w1:      in   std_logic_vector(9 - 1 downto 0);
26     w2:      in   std_logic_vector(9 - 1 downto 0);
27     w3:      in   std_logic_vector(9 - 1 downto 0);
28     w4:      in   std_logic_vector(9 - 1 downto 0);
29     w5:      in   std_logic_vector(9 - 1 downto 0);
30     w6:      in   std_logic_vector(9 - 1 downto 0);
31     w7:      in   std_logic_vector(9 - 1 downto 0);
32     w8:      in   std_logic_vector(9 - 1 downto 0);
33     w9:      in   std_logic_vector(9 - 1 downto 0);
34     -- input bias b, 9 bit, fpv, [-1, 1]
35     b:      in   std_logic_vector(9 - 1 downto 0);
36     -- output y, 16 bit, fpv, [-1, 1]
37     y :      out  std_logic_vector(16 - 1 downto 0)
38     -- ;
39 );
40 end Perceptron;
41
42 -- /**/ /**/ /**/
43 architecture J_perceptron of Perceptron is
44     -- "registers" input
45     signal reg_x0      :   std_logic_vector(8 - 1 downto 0);
46     signal reg_x1      :   std_logic_vector(8 - 1 downto 0);
47     signal reg_x2      :   std_logic_vector(8 - 1 downto 0);
48     signal reg_x3      :   std_logic_vector(8 - 1 downto 0);
49     signal reg_x4      :   std_logic_vector(8 - 1 downto 0);
50     signal reg_x5      :   std_logic_vector(8 - 1 downto 0);
51     signal reg_x6      :   std_logic_vector(8 - 1 downto 0);
52     signal reg_x7      :   std_logic_vector(8 - 1 downto 0);
53     signal reg_x8      :   std_logic_vector(8 - 1 downto 0);
54     signal reg_x9      :   std_logic_vector(8 - 1 downto 0);
55     signal reg_w0      :   std_logic_vector(9 - 1 downto 0);
56     signal reg_w1      :   std_logic_vector(9 - 1 downto 0);
57     signal reg_w2      :   std_logic_vector(9 - 1 downto 0);
58     signal reg_w3      :   std_logic_vector(9 - 1 downto 0);
59     signal reg_w4      :   std_logic_vector(9 - 1 downto 0);
60     signal reg_w5      :   std_logic_vector(9 - 1 downto 0);
61     signal reg_w6      :   std_logic_vector(9 - 1 downto 0);
62     signal reg_w7      :   std_logic_vector(9 - 1 downto 0);
63     signal reg_w8      :   std_logic_vector(9 - 1 downto 0);
64     signal reg_w9      :   std_logic_vector(9 - 1 downto 0);
65     signal reg_bias    :   std_logic_vector(9 - 1 downto 0);
66     -- "register" output
67     -- s reg out :
68     signal reg_y      :   std_logic_vector(16 - 1 downto 0);

```

```

69
70  -- sum-prod of 2int + 6|7 fract -> 17 = 4 int 13 fract in [-1,1]
71  -- sum 11 of these -> 21 = 8 int 13 fract in [-11,11]
72  -- ->>> 5 int 13 fract suffice to avoid OF, -> all in 18 bits 5
    int 13 fract
73      signal esp_bias      :    std_logic_vector(18 - 1 downto 0);
74      signal prod_0        :    std_logic_vector(18 - 1 downto 0);
75      signal prod_1        :    std_logic_vector(18 - 1 downto 0);
76      signal prod_2        :    std_logic_vector(18 - 1 downto 0);
77      signal prod_3        :    std_logic_vector(18 - 1 downto 0);
78      signal prod_4        :    std_logic_vector(18 - 1 downto 0);
79      signal prod_5        :    std_logic_vector(18 - 1 downto 0);
80      signal prod_6        :    std_logic_vector(18 - 1 downto 0);
81      signal prod_7        :    std_logic_vector(18 - 1 downto 0);
82      signal prod_8        :    std_logic_vector(18 - 1 downto 0);
83      signal prod_9        :    std_logic_vector(18 - 1 downto 0);
84
85      signal sum_0to1      :    std_logic_vector(18 - 1 downto 0);
86  -- group 2
87      signal sum_2to3      :    std_logic_vector(18 - 1 downto 0);
88  -- group 2
89      signal sum_4to5      :    std_logic_vector(18 - 1 downto 0);
90  -- group 2
91      signal sum_6to7      :    std_logic_vector(18 - 1 downto 0);
92  -- group 2
93      signal sum_8to9      :    std_logic_vector(18 - 1 downto 0);
94  -- group 2
95      signal sum_0to3      :    std_logic_vector(18 - 1 downto 0);
96  -- group 4
97      signal sum_4to7      :    std_logic_vector(18 - 1 downto 0);
98  -- group 4
99      signal sum_8to9_bs   :    std_logic_vector(18 - 1 downto 0);
100  -- group 3
101      signal sum_0to7      :    std_logic_vector(18 - 1 downto 0);
102  -- group 8
103
104      -- sum_product_result  -- [ "-11" , "+11" ], 5 int 13 fract
105      signal sp_res:        std_logic_vector(18 - 1 downto 0);
106      signal sp_plus_two:   std_logic_vector(18 - 1 downto 0);
107      signal cut_sp_plus_two: std_logic_vector(16 - 1 downto 0);
108
109      -- filter output _y_   -- [ "-1" , "+1" ], 5 int 13 fract
110      signal value_to_out:   std_logic_vector(16 - 1 downto 0);
111      signal af_out:         std_logic_vector(16 - 1 downto 0);
112
113  -- end signals

```

```

105
106 -- begin constants
107 --
108 "0123456789ABCDEF01"2345678"
109 -- 18 bits : 5 int 13 fract :
110 constant DUE_POS_18bit : std_logic_vector(18 - 1 downto 0)
:= "000100000000000000" ;
111 constant DUE_NEG_18bit : std_logic_vector(18 - 1 downto 0)
:= "111100000000000000" ;
112 constant UN_MEZ_18bit : std_logic_vector(18 - 1 downto 0)
:= "000001000000000000" ;
113 -- 16 bits : 3 int 13 fract
114 constant ZERO_16b : std_logic_vector(16 - 1 downto 0)
:= ( others => '0' ) ; -- 0
115 constant PUN0_16b : std_logic_vector(16 - 1 downto 0)
:= ("0010000000000000") ; -- +1
116 --
117 "0123456789ABCDEF"012345678"
118 -- end constants
119
120 -- architecture description [begin, end)
121 -- architecture begin
122 begin --J_perceptron
123     mainBlock: process ( clk,
124                         resetn
125                     )
126     begin
127         -- handle @resetn==0
128         if ( resetn = '0' ) then
129             -- all inputs are considered 0
130             reg_x0 <= (others => '0') ;
131             reg_x1 <= (others => '0') ;
132             reg_x2 <= (others => '0') ;
133             reg_x3 <= (others => '0') ;
134             reg_x4 <= (others => '0') ;
135             reg_x5 <= (others => '0') ;
136             reg_x6 <= (others => '0') ;
137             reg_x7 <= (others => '0') ;
138             reg_x8 <= (others => '0') ;
139             reg_x9 <= (others => '0') ;
140             reg_w0 <= (others => '0') ;
141             reg_w1 <= (others => '0') ;
142             reg_w2 <= (others => '0') ;
143             reg_w3 <= (others => '0') ;
144             reg_w4 <= (others => '0') ;
145             reg_w5 <= (others => '0') ;

```

```

144         reg_w6      <= (others => '0') ;
145         reg_w7      <= (others => '0') ;
146         reg_w8      <= (others => '0') ;
147         reg_w9      <= (others => '0') ;
148         -- bias
149         reg_bias     <= (others => '0') ;
150         -- s reg uscita :
151         -- output default 0
152         reg_y        <= (others => '0') ;
153     elsif ( rising_edge(clk) ) then
154         -- acquire input (8 or 9 bit)
155         reg_x0       <= ( x0 )      ;
156         reg_x1       <= ( x1 )      ;
157         reg_x2       <= ( x2 )      ;
158         reg_x3       <= ( x3 )      ;
159         reg_x4       <= ( x4 )      ;
160         reg_x5       <= ( x5 )      ;
161         reg_x6       <= ( x6 )      ;
162         reg_x7       <= ( x7 )      ;
163         reg_x8       <= ( x8 )      ;
164         reg_x9       <= ( x9 )      ;
165         reg_w0       <= ( w0 )      ;
166         reg_w1       <= ( w1 )      ;
167         reg_w2       <= ( w2 )      ;
168         reg_w3       <= ( w3 )      ;
169         reg_w4       <= ( w4 )      ;
170         reg_w5       <= ( w5 )      ;
171         reg_w6       <= ( w6 )      ;
172         reg_w7       <= ( w7 )      ;
173         reg_w8       <= ( w8 )      ;
174         reg_w9       <= ( w9 )      ;
175         reg_bias     <= ( b )      ;
176         -- s reg uscita :
177         reg_y        <= af_out      ;
178     end if;
179 end process mainBlock;
180
181 -- s reg uscita :
182 y <= reg_y ;
183
184 prod_0      <= std_logic_vector( to_signed ( ( to_integer(
signed(reg_x0)) * to_integer(signed(reg_w0)) ) , 18 ) ) ;
185 prod_1      <= std_logic_vector( to_signed ( ( to_integer(
signed(reg_x1)) * to_integer(signed(reg_w1)) ) , 18 ) ) ;
186 prod_2      <= std_logic_vector( to_signed ( ( to_integer(
signed(reg_x2)) * to_integer(signed(reg_w2)) ) , 18 ) ) ;

```

```

187     prod_3      <= std_logic_vector( to_signed ( ( to_integer(
signed(reg_x3)) * to_integer(signed(reg_w3)) ) , 18 ) ) ;
188     prod_4      <= std_logic_vector( to_signed ( ( to_integer(
signed(reg_x4)) * to_integer(signed(reg_w4)) ) , 18 ) ) ;
189     prod_5      <= std_logic_vector( to_signed ( ( to_integer(
signed(reg_x5)) * to_integer(signed(reg_w5)) ) , 18 ) ) ;
190     prod_6      <= std_logic_vector( to_signed ( ( to_integer(
signed(reg_x6)) * to_integer(signed(reg_w6)) ) , 18 ) ) ;
191     prod_7      <= std_logic_vector( to_signed ( ( to_integer(
signed(reg_x7)) * to_integer(signed(reg_w7)) ) , 18 ) ) ;
192     prod_8      <= std_logic_vector( to_signed ( ( to_integer(
signed(reg_x8)) * to_integer(signed(reg_w8)) ) , 18 ) ) ;
193     prod_9      <= std_logic_vector( to_signed ( ( to_integer(
signed(reg_x9)) * to_integer(signed(reg_w9)) ) , 18 ) ) ;
194
195     -- reg_bias is signed 9 bit, signed expand to 18 bit
196     esp_bias    <= ( reg_bias(9-1) & reg_bias(9-1) & reg_bias
(9-1) & reg_bias(9-1) & reg_bias(9-1) & reg_bias(9-1) & reg_bias
(9-1) & reg_bias(9-1) & reg_bias(9-1) & reg_bias );
197
198     -- from considerations above, 5 int 13 fract enough to not
worry for OF,
199     -- sum in BALANCED BINARY TREE
200     sum_0to1     <= ( prod_0 + prod_1 ) ;      -- group 2
201     sum_2to3     <= ( prod_2 + prod_3 ) ;      -- group 2
202     sum_4to5     <= ( prod_4 + prod_5 ) ;      -- group 2
203     sum_6to7     <= ( prod_6 + prod_7 ) ;      -- group 2
204     sum_8to9     <= ( prod_8 + prod_9 ) ;      -- group 2
205
206     sum_0to3     <= ( sum_0to1 + sum_2to3 ) ; -- group 4
207     sum_4to7     <= ( sum_4to5 + sum_6to7 ) ; -- group 4
208     sum_8to9_bs  <= ( sum_8to9 + esp_bias ) ; -- group 3
209
210     sum_0to7     <= ( sum_0to3 + sum_4to7 ) ; -- group 8
211
212     sp_res    <= ( sum_0to7 + sum_8to9_bs ) ;
213     -- non balanced sum (sequential):
214     -- sp_res <= ( prod_0 + prod_1 + prod_2 + prod_3 + prod_4 +
prod_5 + prod_6 + prod_7 + prod_8 + prod_9 + esp_bias );
215
216     -- !!! OUT_DECISOR !!!
217     -- let sp_res = x
218     -- _ if x > 2, 1
219     -- _ if x <2, 0
220     -- _ else : [ (x+2)/4 ] == [ (x/2) + 1/2 ]
221     -- from given, note : x+2 GE0 && x+2 AE0

```

```

222      -- sp_plus_two :      00XXX.fract_13,
223      -- after shift :      SS00X.XXfract_11
224      -- reduce to 16 bit :  00X.XXfract_11
225      -- i will just cut after addition and all is good
226
227      sp_plus_two <= sp_res + DUE_POS_18bit;
228      cut_sp_plus_two <= ( sp_plus_two((18-1) downto 2 ) );
229
230      -- no reg uscita :
231      -- y <=      PUN0_16b when ( signed(sp_res) > signed(
DUE_POS_18bit) ) else
232      --          ZERO_16b when ( signed(sp_res) < signed(
DUE_NEG_18bit) ) else
233      --          cut_sp_plus_two ;
234
235      -- s reg uscita :
236      af_out <=      PUN0_16b when ( signed(sp_res) > signed(
DUE_POS_18bit) ) else
237      --          ZERO_16b when ( signed(sp_res) < signed(
DUE_NEG_18bit) ) else
238      --          cut_sp_plus_two ;
239
240      -- (end architecture J_perceptron)
241      end J_perceptron;

```

3.3 Test Bench : code

The test bench (TB) structure will be better described in the next section; however, it is important to notice that it has been organized in order to behave like ordinary procedural code. The TB initializes an instance of the Perceptron, then monitors its output to verify that it behaves consistently with the specifications, adjourning the "fail" signal if any problem occurs.

```

1  library IEEE;
2  use IEEE.std_logic_1164.all;
3  use IEEE.numeric_std.all;
4  use IEEE.std_logic_unsigned.all;
5
6  entity Perceptron_tb is
7  end Perceptron_tb ;
8
9  architecture tb of Perceptron_tb is
10     constant CLK_PERIOD : time := 400 ns;
11     constant THREE_QUARTERS_PERIOD : time := 300 ns;
12

```

```

13  constant ZERO_15b      : std_logic_vector(15 - 1 downto 0) := (
others => '0') ;-- 0
14  constant ZERO_6b      : std_logic_vector(6 - 1 downto 0) := (
others => '0') ;-- 0      -- per 7f
15  constant ZERO_7b      : std_logic_vector(7 - 1 downto 0) := (
others => '0') ;-- 0      -- per 6f
16  constant UNO_15b      : std_logic_vector(15 - 1 downto 0) := (
others => '1') ;-- 1
17  constant UNO_6b      : std_logic_vector(6 - 1 downto 0) := (
others => '1') ;-- 1      -- per 7f
18  constant UNO_7b      : std_logic_vector(7 - 1 downto 0) := (
others => '1') ;-- 1      -- per 6f
19
20  constant PUNO_30b_7f  : std_logic_vector(30 - 1 downto 0) := (
ZERO_15b & ZERO_6b & "01000000" ) ; -- +1
21  constant PHLF_30b_7f  : std_logic_vector(30 - 1 downto 0) := (
ZERO_15b & ZERO_6b & "00100000" ) ; -- +0.5
22  constant PQRT_30b_7f  : std_logic_vector(30 - 1 downto 0) := (
ZERO_15b & ZERO_6b & "00010000" ) ; -- +0.25
23  constant NQRT_30b_7f  : std_logic_vector(30 - 1 downto 0) := (
UNO_15b  & UNO_6b  & "11110000" ) ; -- -0.25
24  constant NHLF_30b_7f  : std_logic_vector(30 - 1 downto 0) := (
UNO_15b  & UNO_6b  & "11100000" ) ; -- -0.5
25  constant NUNO_30b_7f  : std_logic_vector(30 - 1 downto 0) := (
UNO_15b  & UNO_6b  & "11000000" ) ; -- -1
26
27  constant PUNO_30b_6f  : std_logic_vector(30 - 1 downto 0) :=
( ZERO_15b & ZERO_7b & "01000000" ) ; -- +1
28  constant PHLF_30b_6f  : std_logic_vector(30 - 1 downto 0) :=
( ZERO_15b & ZERO_7b & "00100000" ) ; -- +0.5
29  constant PQRT_30b_6f  : std_logic_vector(30 - 1 downto 0) :=
( ZERO_15b & ZERO_7b & "00010000" ) ; -- +0.25
30  constant NQRT_30b_6f  : std_logic_vector(30 - 1 downto 0) :=
( UNO_15b  & UNO_7b  & "11110000" ) ; -- -0.25
31  constant NHLF_30b_6f  : std_logic_vector(30 - 1 downto 0) :=
( UNO_15b  & UNO_7b  & "11100000" ) ; -- -0.5
32  constant NUNO_30b_6f  : std_logic_vector(30 - 1 downto 0) :=
( UNO_15b  & UNO_7b  & "11000000" ) ; -- -1
33
34  constant ZERO_30b      : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;-- 0
35
36  -- y ha 13 fract e 3 interi, ed      in [0, 1]
37  -- il +1 sarebbe 001_000_000_000_000_0
38  constant ZERO_16b      : std_logic_vector(16 - 1 downto 0) := (
others => '0') ;-- 0

```



```

39     constant PUN0_16b      : std_logic_vector(16 - 1 downto 0) := (
    "0010000000000000") ; -- +1
40     constant PHLF_16b      : std_logic_vector(16 - 1 downto 0) := (
    "0001000000000000") ; -- 0
41
42
43     component Perceptron
44     port(
45         clk      : in std_logic;
46         resetn    : in std_logic; -- active low
47         -- input xi a 8 bit, in virgola fissa, tra -1 e 1
48         x0        : in std_logic_vector(8 - 1 downto 0);
49         x1        : in std_logic_vector(8 - 1 downto 0);
50         x2        : in std_logic_vector(8 - 1 downto 0);
51         x3        : in std_logic_vector(8 - 1 downto 0);
52         x4        : in std_logic_vector(8 - 1 downto 0);
53         x5        : in std_logic_vector(8 - 1 downto 0);
54         x6        : in std_logic_vector(8 - 1 downto 0);
55         x7        : in std_logic_vector(8 - 1 downto 0);
56         x8        : in std_logic_vector(8 - 1 downto 0);
57         x9        : in std_logic_vector(8 - 1 downto 0);
58         -- input weights wi a 9 bit, in virgola fissa, tra -1 e
1         w0        : in std_logic_vector(9 - 1 downto 0);
59         w1        : in std_logic_vector(9 - 1 downto 0);
60         w2        : in std_logic_vector(9 - 1 downto 0);
61         w3        : in std_logic_vector(9 - 1 downto 0);
62         w4        : in std_logic_vector(9 - 1 downto 0);
63         w5        : in std_logic_vector(9 - 1 downto 0);
64         w6        : in std_logic_vector(9 - 1 downto 0);
65         w7        : in std_logic_vector(9 - 1 downto 0);
66         w8        : in std_logic_vector(9 - 1 downto 0);
67         w9        : in std_logic_vector(9 - 1 downto 0);
68         -- input bias a 9 bit, in virgola fissa, tra -1 e 1
69         b          : in std_logic_vector(9 - 1 downto 0);
70         -- output [-1, 1]
71         y          : out std_logic_vector(16 - 1 downto 0)
72     );
73
74     end component;
75
76     -- control signals
77     -- timings signals
78     signal clk_tb      : std_logic      := '0';
79     signal resetn_tb   : std_logic      := '0';
80     -- test estimators
81     signal success      : std_logic      := '0';

```

```

82     signal fail      : std_logic      := '0';
83     signal pass_input: std_logic      := '0';
84     signal testing   : boolean := true;
85     -- mi preparo tutti i registri a 0
86     -- input vals
87     signal tb_reg_x0 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
88     signal tb_reg_x1 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
89     signal tb_reg_x2 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
90     signal tb_reg_x3 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
91     signal tb_reg_x4 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
92     signal tb_reg_x5 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
93     signal tb_reg_x6 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
94     signal tb_reg_x7 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
95     signal tb_reg_x8 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
96     signal tb_reg_x9 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
97     -- input weights
98     signal tb_reg_w0 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
99     signal tb_reg_w1 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
100    signal tb_reg_w2 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
101    signal tb_reg_w3 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
102    signal tb_reg_w4 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
103    signal tb_reg_w5 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
104    signal tb_reg_w6 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
105    signal tb_reg_w7 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
106    signal tb_reg_w8 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
107    signal tb_reg_w9 : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;

```

```

108  -- input system bias
109  signal tb_reg_bias : std_logic_vector(30 - 1 downto 0) := (
others => '0' ) ;
110  -- output filtered result
111  signal tbs_y_rcv : std_logic_vector(16 - 1 downto 0) := (
others => '0' ) ;
112  -- end signals
113
114  -- begin testing
115  begin
116  clk_tb <= not clk_tb after CLK_PERIOD / 2 when testing else
'0';
117  -- attach device to test (device under test)
118  dut : Perceptron
119  port map(
120      clk          => clk_tb,
121      resetn       => resetn_tb,
122      x0           => tb_reg_x0( 8-1 downto 0 ) ,
123      x1           => tb_reg_x1( 8-1 downto 0 ) ,
124      x2           => tb_reg_x2( 8-1 downto 0 ) ,
125      x3           => tb_reg_x3( 8-1 downto 0 ) ,
126      x4           => tb_reg_x4( 8-1 downto 0 ) ,
127      x5           => tb_reg_x5( 8-1 downto 0 ) ,
128      x6           => tb_reg_x6( 8-1 downto 0 ) ,
129      x7           => tb_reg_x7( 8-1 downto 0 ) ,
130      x8           => tb_reg_x8( 8-1 downto 0 ) ,
131      x9           => tb_reg_x9( 8-1 downto 0 ) ,
132      w0           => tb_reg_w0( 9-1 downto 0 ) ,
133      w1           => tb_reg_w1( 9-1 downto 0 ) ,
134      w2           => tb_reg_w2( 9-1 downto 0 ) ,
135      w3           => tb_reg_w3( 9-1 downto 0 ) ,
136      w4           => tb_reg_w4( 9-1 downto 0 ) ,
137      w5           => tb_reg_w5( 9-1 downto 0 ) ,
138      w6           => tb_reg_w6( 9-1 downto 0 ) ,
139      w7           => tb_reg_w7( 9-1 downto 0 ) ,
140      w8           => tb_reg_w8( 9-1 downto 0 ) ,
141      w9           => tb_reg_w9( 9-1 downto 0 ) ,
142      b            => tb_reg_bias( 9-1 downto 0 ) ,
143      y            => tbs_y_rcv
144  );
145
146  -- testing "program"
147  TB_PROC: process -- (clk_tb)
148  begin
149      if testing then
150          wait until rising_edge(clk_tb);

```

```

151         resetn_tb <= '0' ;
152         wait until rising_edge(clk_tb);
153         -- end reset (reset begins at 0 in tb, now i put it
at 1)
154         resetn_tb <= '1' ;
155         -- give initial data: weights not 0, x all 0
156         wait until rising_edge(clk_tb);
157         -- x ha 6 bit fract e 2 interi, metto tutto a 0 _ il
reg      a 30 bit
158         tb_reg_x0 <= ZERO_30b ;
159         tb_reg_x1 <= ZERO_30b ;
160         tb_reg_x2 <= ZERO_30b ;
161         tb_reg_x3 <= ZERO_30b ;
162         tb_reg_x4 <= ZERO_30b ;
163         tb_reg_x5 <= ZERO_30b ;
164         tb_reg_x6 <= ZERO_30b ;
165         tb_reg_x7 <= ZERO_30b ;
166         tb_reg_x8 <= ZERO_30b ;
167         tb_reg_x9 <= ZERO_30b ;
168         -- w e bias hanno 7 bit fract e 2 interi _ il reg
a 30 bit
169         tb_reg_w0 <= PUN0_30b_7f ;    -- +1
170         tb_reg_w1 <= PUN0_30b_7f ;    -- +1
171         tb_reg_w2 <= PUN0_30b_7f ;    -- +1
172         tb_reg_w3 <= PHLF_30b_7f ;    -- +0.5
173         tb_reg_w4 <= PQRT_30b_7f ;    -- +0.25
174         tb_reg_w5 <= NUN0_30b_7f ;    -- -1
175         tb_reg_w6 <= NUN0_30b_7f ;    -- -1
176         tb_reg_w7 <= NUN0_30b_7f ;    -- -1
177         tb_reg_w8 <= NHLF_30b_7f ;    -- -0.5
178         tb_reg_w9 <= NQRT_30b_7f ;    -- -0.25
179         -- bias come sopra
180         tb_reg_bias <= ZERO_30b ;      -- +0
181         pass_input <= '1';
182         -- tutti gli input sono 0, output dovrebbe essere 1/2;
183         wait until rising_edge(clk_tb);
184         pass_input <= '0';
185         wait until rising_edge(clk_tb);
186         wait for 15 ns;
187         if tbs_y_rcv /= PHLF_16b then
188             fail<= '1' ;
189         end if;
190         -- daccapino
191         -- test x values:
192         -- pos -> 1
193         wait until rising_edge(clk_tb);

```

```

194     pass_input <= '1';
195     wait for THREE_QUARTERS_PERIOD ;
196     -- cambio solo un paio di registri intanto
197     tb_reg_x0 <= PUN0_30b_6f      ; -- * +1
198     tb_reg_x1 <= PUN0_30b_6f      ; -- * +1
199     tb_reg_x2 <= PUN0_30b_6f      ; -- * +1
200     tb_reg_x5 <= PUN0_30b_6f      ; -- * -1
201     tb_reg_x6 <= ZERO_30b         ; -- * -1
202     tb_reg_x7 <= ZERO_30b         ; -- * -1
203     -- adesso y dovrebbe essere 1+1+1-1-0-0 = 2 quindi 1
204     wait until rising_edge(clk_tb);
205     pass_input <= '0';
206     wait until rising_edge(clk_tb);
207     wait for 15 ns;
208     if tbs_y_rcv /= PUN0_16b then
209         fail<= '1' ;
210     end if;
211     -- neg -> 0
212     wait until rising_edge(clk_tb);
213     pass_input <= '1';
214     wait for THREE_QUARTERS_PERIOD ;
215     -- cambio solo un paio di registri intanto
216     tb_reg_x0 <= NUN0_30b_6f      ; -- * +1
217     tb_reg_x1 <= NUN0_30b_6f      ; -- * +1
218     tb_reg_x2 <= NUN0_30b_6f      ; -- * +1
219     tb_reg_x5 <= NUN0_30b_6f      ; -- * -1
220     tb_reg_x6 <= ZERO_30b         ; -- * -1
221     tb_reg_x7 <= ZERO_30b         ; -- * -1
222     -- adesso y dovrebbe essere -1 -1 -1 +1 +0 +0 = -2
    quindi 0
223     wait until rising_edge(clk_tb);
224     pass_input <= '0';
225     wait until rising_edge(clk_tb);
226     wait for 15 ns;
227     if tbs_y_rcv /= ZERO_16b then
228         fail<= '1' ;
229     end if;
230     -- zero -> 0.5
231     wait until rising_edge(clk_tb);
232     pass_input <= '1';
233     wait for THREE_QUARTERS_PERIOD ;
234     -- cambio solo un paio di registri intanto
235     tb_reg_x0 <= PUN0_30b_6f      ; -- * +1
236     tb_reg_x1 <= PUN0_30b_6f      ; -- * +1
237     tb_reg_x2 <= PUN0_30b_6f      ; -- * +1
238     tb_reg_x5 <= PUN0_30b_6f      ; -- * -1

```

```

239         tb_reg_x6 <= PUN0_30b_6f      ; -- * -1
240         tb_reg_x7 <= PUN0_30b_6f      ; -- * -1
241         -- adesso y dovrebbe essere -1 -1 -1 +1 +0 +0 = 0
quindi 1/2
242         wait until rising_edge(clk_tb);
243         pass_input <= '0';
244         wait until rising_edge(clk_tb);
245         wait for 15 ns;
246         if tbs_y_rcv /= PHLF_16b then
247             fail<= '1' ;
248         end if;
249     -- daccapone
250     -- test x values:
251     -- pos -> 1
252         wait until rising_edge(clk_tb);
253         pass_input <= '1';
254         wait for THREE_QUARTERS_PERIOD ;
255         -- cambio solo un paio di registri intanto
256         tb_reg_x0 <= PUN0_30b_6f      ; -- * +1
257         tb_reg_x1 <= PUN0_30b_6f      ; -- * +1
258         tb_reg_x2 <= PUN0_30b_6f      ; -- * +1
259         tb_reg_x5 <= PUN0_30b_6f      ; -- * -1
260         tb_reg_x6 <= ZERO_30b         ; -- * -1
261         tb_reg_x7 <= ZERO_30b         ; -- * -1
262         -- adesso y dovrebbe essere 1+1+1-1-0-0 = 2 quindi 1
263         wait until rising_edge(clk_tb);
264         pass_input <= '0';
265         wait until rising_edge(clk_tb);
266         wait for 15 ns;
267         if tbs_y_rcv /= PUN0_16b then
268             fail<= '1' ;
269         end if;
270     -- neg -> 0
271         wait until rising_edge(clk_tb);
272         pass_input <= '1';
273         wait for THREE_QUARTERS_PERIOD ;
274         -- cambio solo un paio di registri intanto
275         tb_reg_x0 <= NUN0_30b_6f      ; -- * +1
276         tb_reg_x1 <= NUN0_30b_6f      ; -- * +1
277         tb_reg_x2 <= NUN0_30b_6f      ; -- * +1
278         tb_reg_x5 <= NUN0_30b_6f      ; -- * -1
279         tb_reg_x6 <= ZERO_30b         ; -- * -1
280         tb_reg_x7 <= ZERO_30b         ; -- * -1
281         -- adesso y dovrebbe essere -1 -1 -1 +1 +0 +0 = -2
quindi 0
282         wait until rising_edge(clk_tb);

```

```

283     pass_input <= '0';
284     wait until rising_edge(clk_tb);
285     wait for 15 ns;
286     if tbs_y_rcv /= ZERO_16b then
287         fail<= '1' ;
288     end if;
289     -- zero -> 0.5
290     wait until rising_edge(clk_tb);
291     pass_input <= '1';
292     wait for THREE_QUARTERS_PERIOD ;
293     -- cambio solo un paio di registri intanto
294     tb_reg_x0 <= PUN0_30b_6f      ; -- * +1
295     tb_reg_x1 <= PUN0_30b_6f      ; -- * +1
296     tb_reg_x2 <= PUN0_30b_6f      ; -- * +1
297     tb_reg_x5 <= PUN0_30b_6f      ; -- * -1
298     tb_reg_x6 <= PUN0_30b_6f      ; -- * -1
299     tb_reg_x7 <= PUN0_30b_6f      ; -- * -1
300     -- adesso y dovrebbe essere -1 -1 -1 +1 +0 +0 = 0
    quindi 1/2
301     wait until rising_edge(clk_tb);
302     pass_input <= '0';
303     wait until rising_edge(clk_tb);
304     wait for 15 ns;
305     if tbs_y_rcv /= PHLF_16b then
306         fail<= '1' ;
307     end if;
308     -- other tests omitted for brevity of print
309     -- ...
310     wait until rising_edge(clk_tb);
311     pass_input <= '1';
312     wait until rising_edge(clk_tb);
313     pass_input <= '0';
314     wait until rising_edge(clk_tb);
315     pass_input <= '1';
316     wait until rising_edge(clk_tb);
317     pass_input <= '0';
318     testing <= false;
319     end if; -- (di testing == 1)
320
321     end process;
322
323     end architecture;
324 -- fine

```

4 Verification and Testing

The above test-bench was prepared to verify the design and the synthesis of the circuit. It can be interpreted as an ordinary procedural script that sets the input parameter of the Perceptron, and then ensures that the device under test (DUT) behaves as specifications require.

The inputs are all assigned at reset so that x_i are all zero, while the weights w_i are assigned and, in the above tests, never changed.

Due to the lengthy writing of the inputs, during the tests only part of them are changed (time by time).

Some testing signals have been added: fail (defaults 0, 1 if any test fails), pass_input (to explicit the beginning of the iteration, falls as stabile inputs are read by Perceptron), testing (to run the clock only during tests).

The iterations after reset are as follows:

1. @*rising_edge*(clk)[i] set pass_input=1 ,
2. after some time (delay in input generation) assign inputs,
3. @*rising_edge*(clk)[i+1] set pass_input=0 , (inputs are received in Perceptron)
4. wait random propagation delay (not needed, also works without this)
5. verify output correctness (all before *rising_edge*(clk)[i+2])

Since the Perceptron receives the inputs at *rising_edge*(clk)[i+1] and is expected to (and does) give correct output at *rising_edge*(clk)[i+2], the elaboration time is within one-clock-period.

It needs to be stated that nothing stops us from using step 4 & 5 to give new inputs, since, as stated, the whole Perceptron works withing a 1-clock-period.

In the above testbench it was not done just to ease readability of the ModelSim output.

The obtained results are coherent with the specifications, so the testing phase has been successfully passed.

4.1 Test-Bench output example

The following image is an example of the TB output obtained with the application "Intel FPGA Starter Edition". As evidenced, the test was successfull (note that "fail" remains 0 all the time).

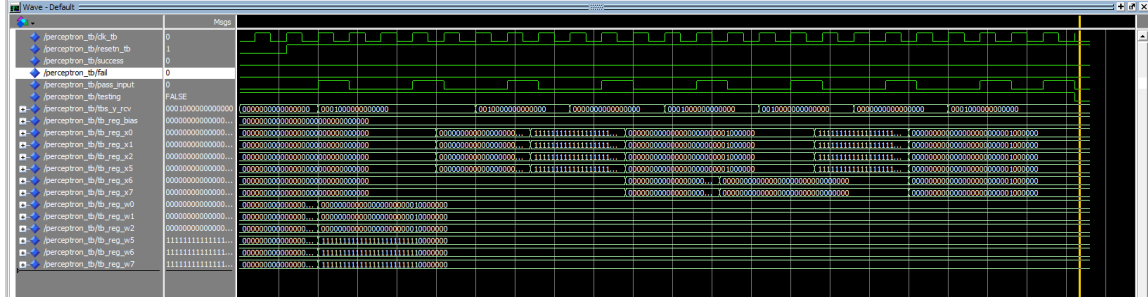


Figure 5: Perceptron and TB signals showed by Intel FPGA Starter Edition.

4.2 Expanding the Test-Bench

Obviously, the testing can be expanded at will, as it was during the writing phase.

The current testing aims to verify the expected output given some "easy to generate" inputs, and, considering the mathematical generality of the algorithm I have implemented, no surprises are to be expected should the inputs vary in the defined input space of $[-1, +1]$.

The above testbench already covers reset, value greater than 2, value less than -2, and value in $[-2, 2]$, so all expected outputs are tested at the least twice (every time successfully).

5 Synthesis and Implementation

After the verification phase, the circuit has been synthesized and implemented using the Vivado Tool for the Zybo Zync-7010 board and the produced results have been studied. VIVADO has been instructed to set the clock speed to 62.5MHz (corresponding to 16ns clock period), accordingly with the board's characteristics.

5.1 Vivado design flow

The following steps of the design flow have been performed using Xilinx Vivado software, which allows to perform RTL Elaboration, Synthesis and Implementation on FPGA and also allows to set constraints and get power or timing reports. To ensure no signal path is omitted during the analysis from the software, all system routes have to be associated with a register-logic-register path, therefore all combinatory logic has to be wrapped between a pair of registers.

This was already achieved while designing the Perceptron : every input and output passes through registers (except clock and resetn).

5.2 RTL

VIVADO produced a logic network made of:

1. 220 cells (i.e. multiplexers, registers, AND gates, etc...);
2. 197 IO ports: ($10 \cdot 8 + 11 \cdot 9$ input bits, 16 output bits, clock and resetn);
3. 778 nets: used to interconnect all other components.

5.3 RTL Elaboration

The RTL Elaboration generated results consistent with the expected structure of the system :

1. registers for the inputs and output;
2. after the MULTipliers, ADDers are disposed in a balanced binary tree;
3. two multiplexers are used to implement the out_decisor.

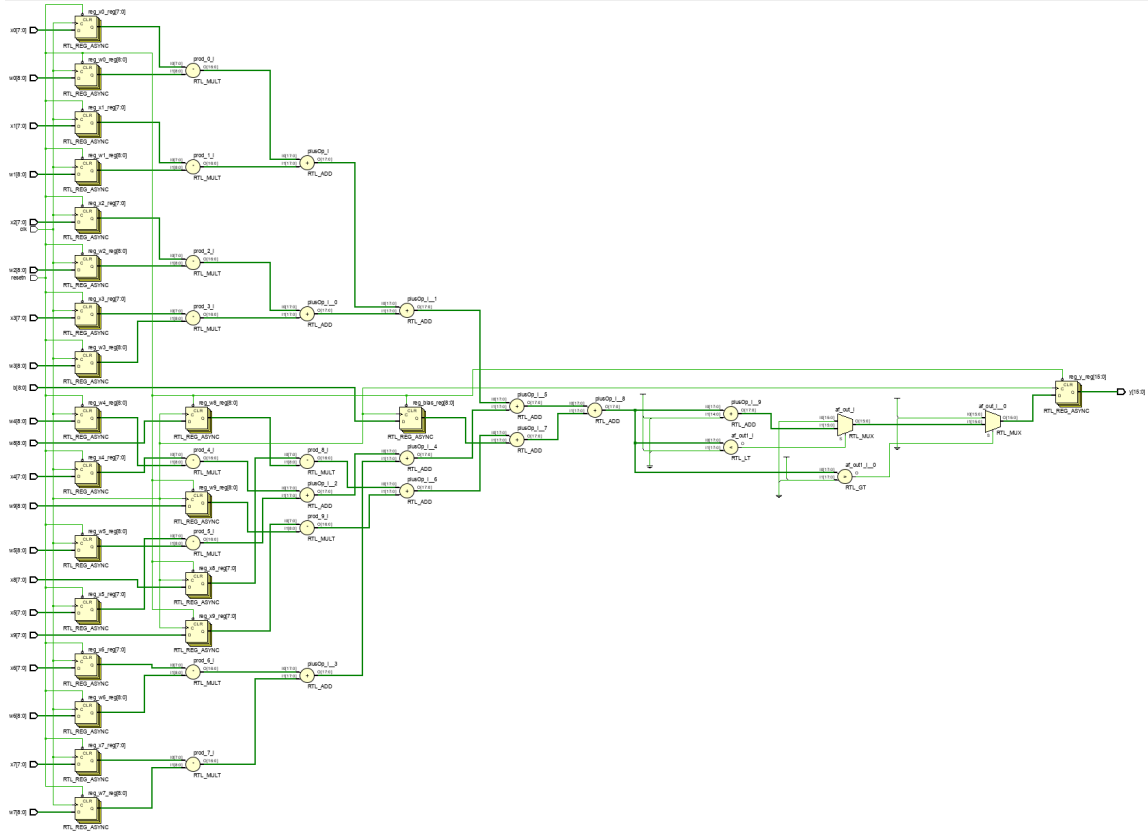


Figure 6: Elaborated RTL Design

5.4 Synthesis and Implementation

The Synthesis runs successfully, but it generates a warning, due to the very high number of inputs, which is unavoidable despite enabling or disabling hierarchy in synthesis.

The Timing Constraints are verified against a **62.5MHz clock (16 ns period)**.

The Implementation yields successful results, although the synthesis is run in *out of context (OOC)* mode, since no mapping of I/O pins is required.

This raises a few extra warnings due to the difficulty in calculating latency on input signals for the board, but accurate timing and power reports can still be obtained.

5.5 Troubleshooting and warnings

No errors were encountered during synthesis and implementation, but a few warnings are found:

- [Netlist 29-101] Netlist 'Perceptron' is not ideal for floorplanning, since the cellview 'Perceptron' contains a large number of primitives. Please consider enabling hierarchy in synthesis if you want to do floorplanning. : This is due to the huge number of inputs of this design, and appears to be unavoidable (even when following documentation, and forum suggestions);
- [Timing 38-242] The property HD.CLK_SRC of clock port "clk" is not set. In out-of-context mode, this prevents timing estimation for clock delay/skew & [Route 35-197] Clock port "clk" does not have an associated HD.CLK_SRC. Without this constraint, timing analysis may not be accurate and upstream checks cannot be done to ensure correct clock placement. : another common warning, it is related to the impossibility to perfectly collect latency data in OOC implementation mode. However, it is possible to easily fix this by setting the requested property to the clock port. To do so, one can simply access the Vivado command line and issue the command: `set_property HD.CLK_SRC BUFGCTRL_X0Y0 [get_ports pclk_in]` which connects a global clock buffer to the clock port of the system.
- [Route 35-198] Port [...] does not have an associated HD.PARTPIN_LOCS, which will prevent the partial routing of the signal [...]. Without this partial route, timing analysis to/from this port will not be accurate, and no routing information for this port can be exported. : this is a warning associated with OOC mode as well and it is similar to the one encountered for clock port, due to the inability to connect external ports to I/O pins. However the warning is not critical and global buffers are only available for clock signals, therefore it has been ignored.

6 Vivado Results

6.1 Critical Path

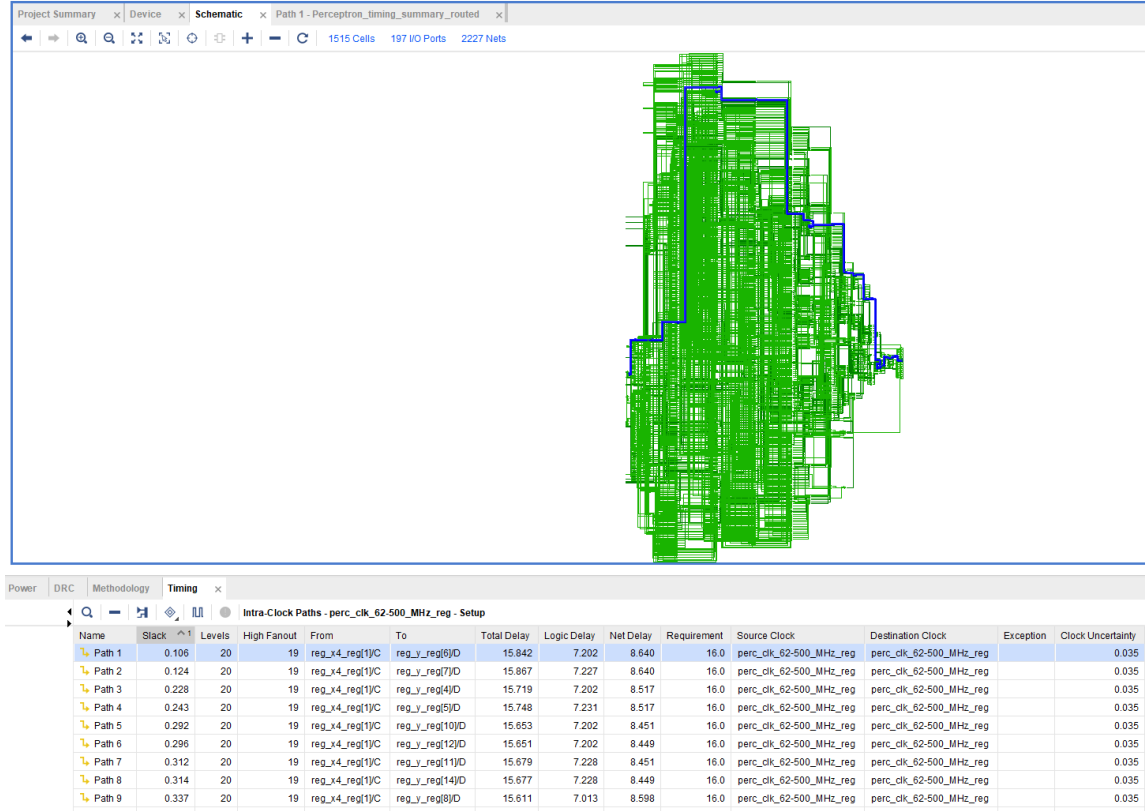


Figure 7: Most Critical Path: Schematic & Timing

The most Critical Paths in both the synthesis and implementation go from the input, through the multipliers and adders, to the multiplexers and output. Since basically all inputs make the same theoretic paths, there is nothing too noteworthy in this behaviour.

Methodology Timing x													
Intra-Clock Paths - perc_clk_62-500_MHz_reg - Setup													
Item	Slack	Levels	High Fanout	From	To	Total Delay	Logic Delay	Net Delay	Requirement	Source Clock	Destination Clock	Exception	Clock Uncertainty
Path 1	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 2	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 3	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 4	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 5	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 6	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 7	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 8	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 9	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 10	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 11	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 12	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 13	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 14	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 15	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 16	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 17	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 18	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 19	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 20	0.106	20	19	reg_x4_reg[1]C	reg_y_reg[6]D	15.842	7.202	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 21	0.124	20	19	reg_x4_reg[1]C	reg_y_reg[7]D	15.867	7.227	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 22	0.124	20	19	reg_x4_reg[1]C	reg_y_reg[7]D	15.867	7.227	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 23	0.124	20	19	reg_x4_reg[1]C	reg_y_reg[7]D	15.867	7.227	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 24	0.124	20	19	reg_x4_reg[1]C	reg_y_reg[7]D	15.867	7.227	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 25	0.124	20	19	reg_x4_reg[1]C	reg_y_reg[7]D	15.867	7.227	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 26	0.124	20	19	reg_x4_reg[1]C	reg_y_reg[7]D	15.867	7.227	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 27	0.124	20	19	reg_x4_reg[1]C	reg_y_reg[7]D	15.867	7.227	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 28	0.124	20	19	reg_x4_reg[1]C	reg_y_reg[7]D	15.867	7.227	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 29	0.124	20	19	reg_x4_reg[1]C	reg_y_reg[7]D	15.867	7.227	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035
Path 30	0.124	20	19	reg_x4_reg[1]C	reg_y_reg[7]D	15.867	7.227	8.640	16.000	perc_clk_62-500_MHz_reg	perc_clk_62-500_MHz_reg		0.035

Figure 8: Timing of the most critical paths

6.2 Timing Report

Tcl Console Messages Log Reports Design Runs Power DRC Methodology Timing x													
Design Timing Summary													
General Information													
Timing Settings													
Clock Summary (1)													
Check Timing (196)													
Intra-Clock Paths													
perc_clk_62-500_MHz_reg													
Setup 0.106 ns (10)													
Hold 1.161 ns (10)													
Pulse Width 7.500 ns (30)													
Inter-Clock Paths													
Other Path Groups													
User Ignored Paths													
Unconstrained Paths													
Timing Summary - impl_1 (saved)													

Setup	Hold	Pulse Width
Worst Negative Slack (WNS): 0.106 ns	Worst Hold Slack (WHS): 1.161 ns	Worst Pulse Width Slack (WPWS): 7.500 ns
Total Negative Slack (TNS): 0.000 ns	Total Hold Slack (THS): 0.000 ns	Total Pulse Width Negative Slack (TPWS): 0.000 ns
Number of Failing Endpoints: 0	Number of Failing Endpoints: 0	Number of Failing Endpoints: 0
Total Number of Endpoints: 16	Total Number of Endpoints: 16	Total Number of Endpoints: 195

All user specified timing constraints are met.

Figure 9: Timing Report

The Timing Report shows quite close margin on the Worst Negative Slack (WNS), and indicates that the **minimum clock-period 16 ns**, corresponding to a **maximum clock-frequency 62.500 MHz**, is extremely close to the best usable.

This frequency, having very small timing tolerance (0.106 ns), might cause some problems in case of clock-jitter.

6.3 Utilization Report

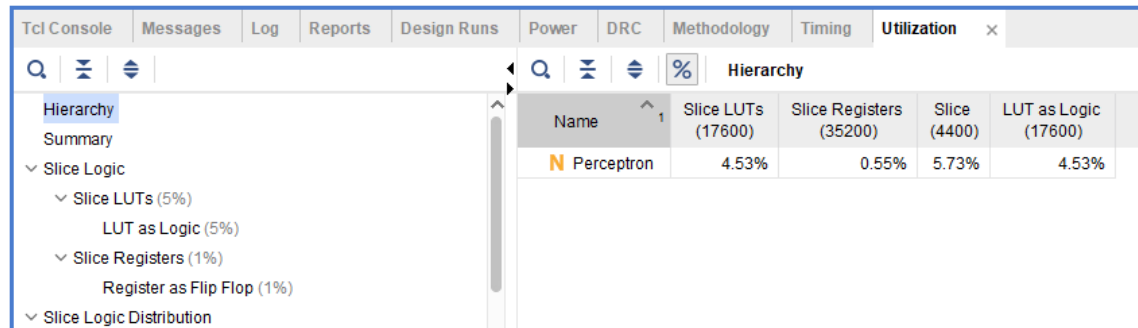


Figure 10: Utilization Report

The Utilization Report indicates a FPGA utilization of 5.73% .

6.4 Power Report

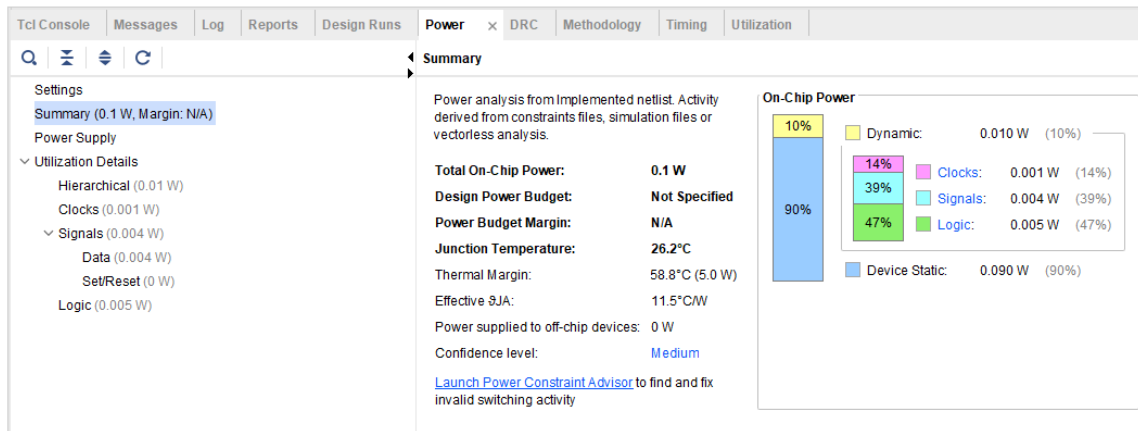


Figure 11: Power Report

The Power Report calculates 0.1W of static power consumption with medium confidence level, mostly related to Logic (47%) and Signals (39%), and an average chip temperature slightly over room temperature (26.2° C) and a wide thermal margin.

Unfortunately, dynamic power consumption is not easy to extract without measuring the system behaviour while it is actively functioning.

7 Final Considerations

The studied circuit is conceptually extremely simple, and definitely clear and concise in what it has to do. The precise -and renowned- application in the field of AI and Deep Learning makes it a very interesting component to study, optimize, build and use.

It can be integrated, with proper timing, into more complex systems of neuron multipliers.

7.1 VIVADO results analysis

Considering the utilization report, it would seem that the circuit is not well suited to be hosted by itself in the considered FPGA (most of the available resources are not needed), but it also indicates that it is designed to perform using quite a simple logic.

More combined or complex implementations will need to deal with the very thin time-margin, but could be optimized (in case of sequential blocks) to reduce the number of registers in separation between "first-input" and "last-output", obviously after proper time analysis of hazards (alee) and worst-paths.